

Improving the performance of nearest neighbor

Recall: nearest neighbor on MNIST

- Images of handwritten digits, represented as vectors in \mathbb{R}^{784} .
- Labels 0 – 9
- Training set: 60,000 points; test set: 10,000 points

Test error of nearest neighbor using Euclidean distance: 3.09%

Examples of errors:

Query



NN



Ideas for improvement: (1) k -NN (2) better distance function.

K-nearest neighbor classification

To classify a new point:

- Find the k nearest neighbors in the training set.
- Return the most common label amongst them.

if $k=1$, we choose 1 NN
 $k=3$, we choose 3 NN

MNIST:

k	1	3	5	7	9	11
Test error (%)	3.09	2.94	3.13	3.10	3.43	3.34

In real life, there's no test set. How to decide which k is best?

Cross-validation

How to estimate the error of k -NN for a particular k ?

10-fold cross-validation


- Divide the training set into 10 equal pieces.
Training set (call it S): 60,000 points
Call the pieces S_1, S_2, \dots, S_{10} : 6,000 points each.
- For each piece S_i :
 - Classify each point in S_i using k -NN with training set $S - S_i$
 - Let ϵ_i = fraction of S_i that is incorrectly classified
- Take the average of these 10 numbers:

10 equal pieces, 9 pieces as training, 1 piece as test

$$\text{estimated error with } k\text{-NN} = \frac{\epsilon_1 + \dots + \epsilon_{10}}{10}$$

Another improvement: better distance functions

The Euclidean (ℓ_2) distance between these two images is very high!



Much better idea: distance measures that are invariant under:

- Small translations and rotations. e.g. *tangent distance*. ✓
- A broader family of natural deformations. e.g. *shape context*. ✓


Test error rates:

ℓ_2	tangent distance	shape context
3.09	1.10	0.63

Related problem: feature selection

Feature selection/reweighting is part of picking a distance function.

And, one noisy feature can wreak havoc with nearest neighbor!



Algorithmic issue: speeding up NN search

Naive search takes time $O(n)$ for training set of size n : slow!

Luckily there are data structures for speeding up nearest neighbor search, like:

- ① Locality sensitive hashing ✓
- ② Ball trees ✓
- ③ K-d trees ✓ ✓ ✓ ✓

These are part of standard Python libraries for NN, and help a lot.

Nearest Neighbor for classification — Build model from data
simplest & most flexible — Use model to make predictions